TOMORROW HAS QUESTIONS

Flying cars. Clever machines. And the wonderfully difficult business of being human.

A field trip through the ideas in Moonshots

With Peter Diamandis, Anousheh Ansari, Salim Ismail and Emad Mostaque

Based on the supplied conversation. Evidence checked 20 October 2026.

[1] Moonshots conversation

Your robot starts Monday

It can take one thing off your hands. What goes first?

A

The chores

Laundry, dishes,
endless tiny errands.

B

The work

Reports, schedules,
repetitive decisions.

C

The hard thinking

Research, calculations,
problems that resist you.

And what would you do with the time?

A show-of-hands question. There is no correct answer.

Four questions for the trip

01 **What can machines do?**

Cars, robots and scientific work

02 **What do we reward?**

XPRIZE and the art of useful incentives

03 **Who chooses the goal?**

Alignment, institutions and shared benefits

04 **What will we choose?**

Purpose, play and life with other people

The episode supplies the questions. The diagrams and audience exercises develop them.

[1] Moonshots conversation

A FLYING CAR IS A WONDERFUL QUESTION

The next question:
who insures the landing?

Roadster discussion, 07:07-16:54. Conceptual illustration. See the evidence notes for qualifications.

[1] Moonshots conversation [11] Tesla Roadster

A demo still has to pass everyday life

1

Capability

Can it do the task?

2

Reliability

Can it repeat it safely?

3

Economics

Does the whole system pay?

4

Institutions

Who permits it and answers for it?

The spectacular moment is only one part of deployment.

Discussion framework inspired by the episode, 12:50-15:56 and 24:37-27:00.

[1] Moonshots conversation

Cheaper thinking changes the economics

A published robotics comparison, with matched caching assumptions

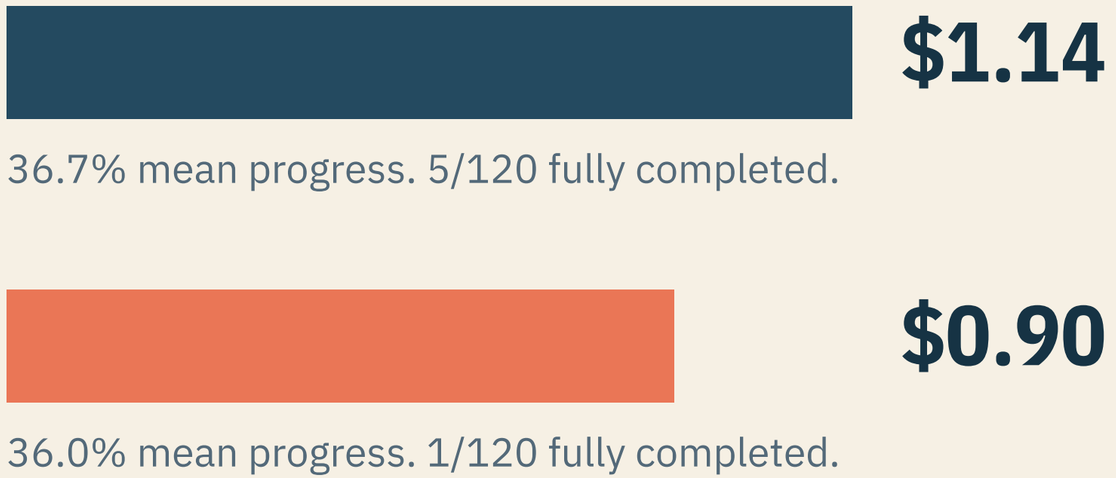

21%

lower estimated model cost per trial

Cheaper attempts still need better completion.

Six selected real-world tasks. Progress is partial task progress, not success rate. Model cost excludes the robot and operations.

[2] Robocurve, RoboDojo-RC Tier 1

A great step can still make a lousy morning

Suppose each step succeeds 99% of the time.

99%

1 step

82%

20 steps

37%

100 steps

Long routines need recovery, checking and safe ways to stop.

Illustrative mathematics, not a robot benchmark. All steps must succeed, with independent failures: 0.99 raised to the number of steps.

YOUR HOME IS THE FINAL BOSS

A sock. A spill.
A very convincing teddy bear.

Useful robots need judgment
when the script runs out.

A household full of exceptions. Episode, 24:37-26:28. Conceptual illustration.

[1] Moonshots conversation

The right body depends on the job

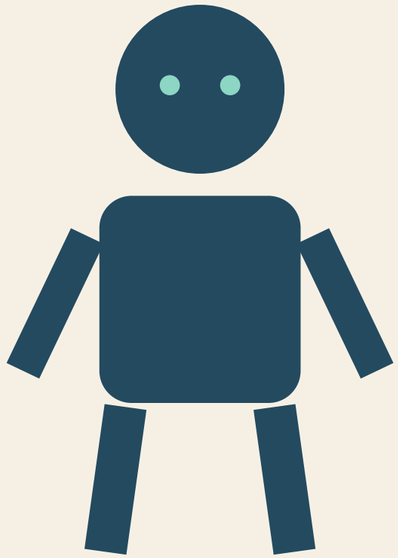

Human-shaped

Works with tools and spaces built around people.

Task-shaped

Trades generality for a simpler, more efficient job.

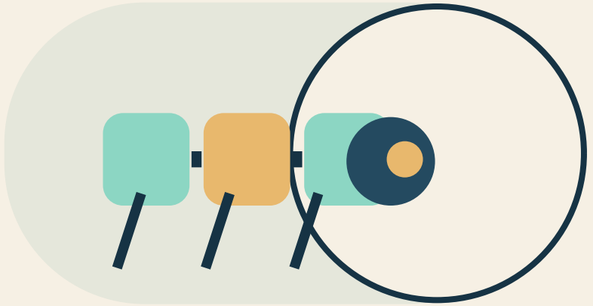

Modular

Changes tools or parts as the job changes.

These are design choices, not a forecast that one form factor will win. Episode, 21:51-28:18.

[1] Moonshots conversation

Physics gets a very fast collaborator

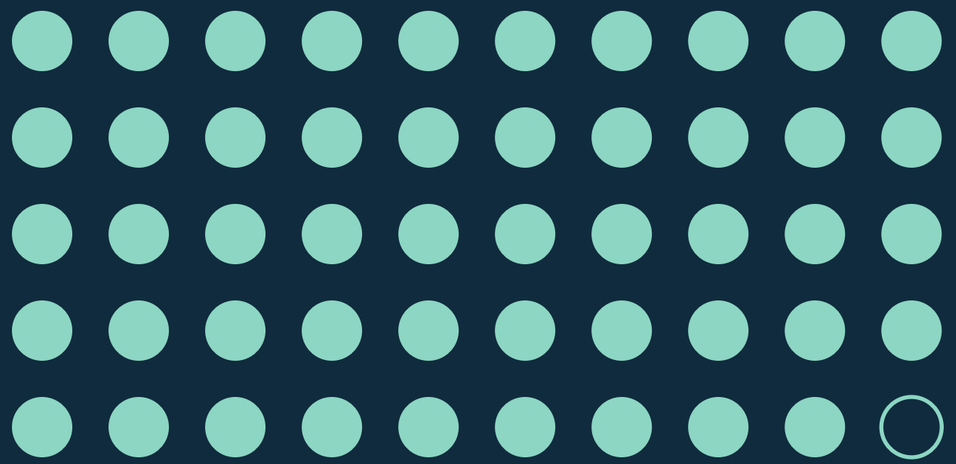

49 / 50

unpublished theoretical physics problems

Under $3 per problem on average

With literature search and computational tools

Choosing an important new question remains a different challenge.

TheoPhysEval authors report this result for GPT-6 Astra in Codex. This preprint measures concrete calculations, not all of scientific discovery.

[3] TheoPhysEval preprint

A PRIZE IS A QUESTION WITH A FINISH LINE

People you never expected may bring the answer.

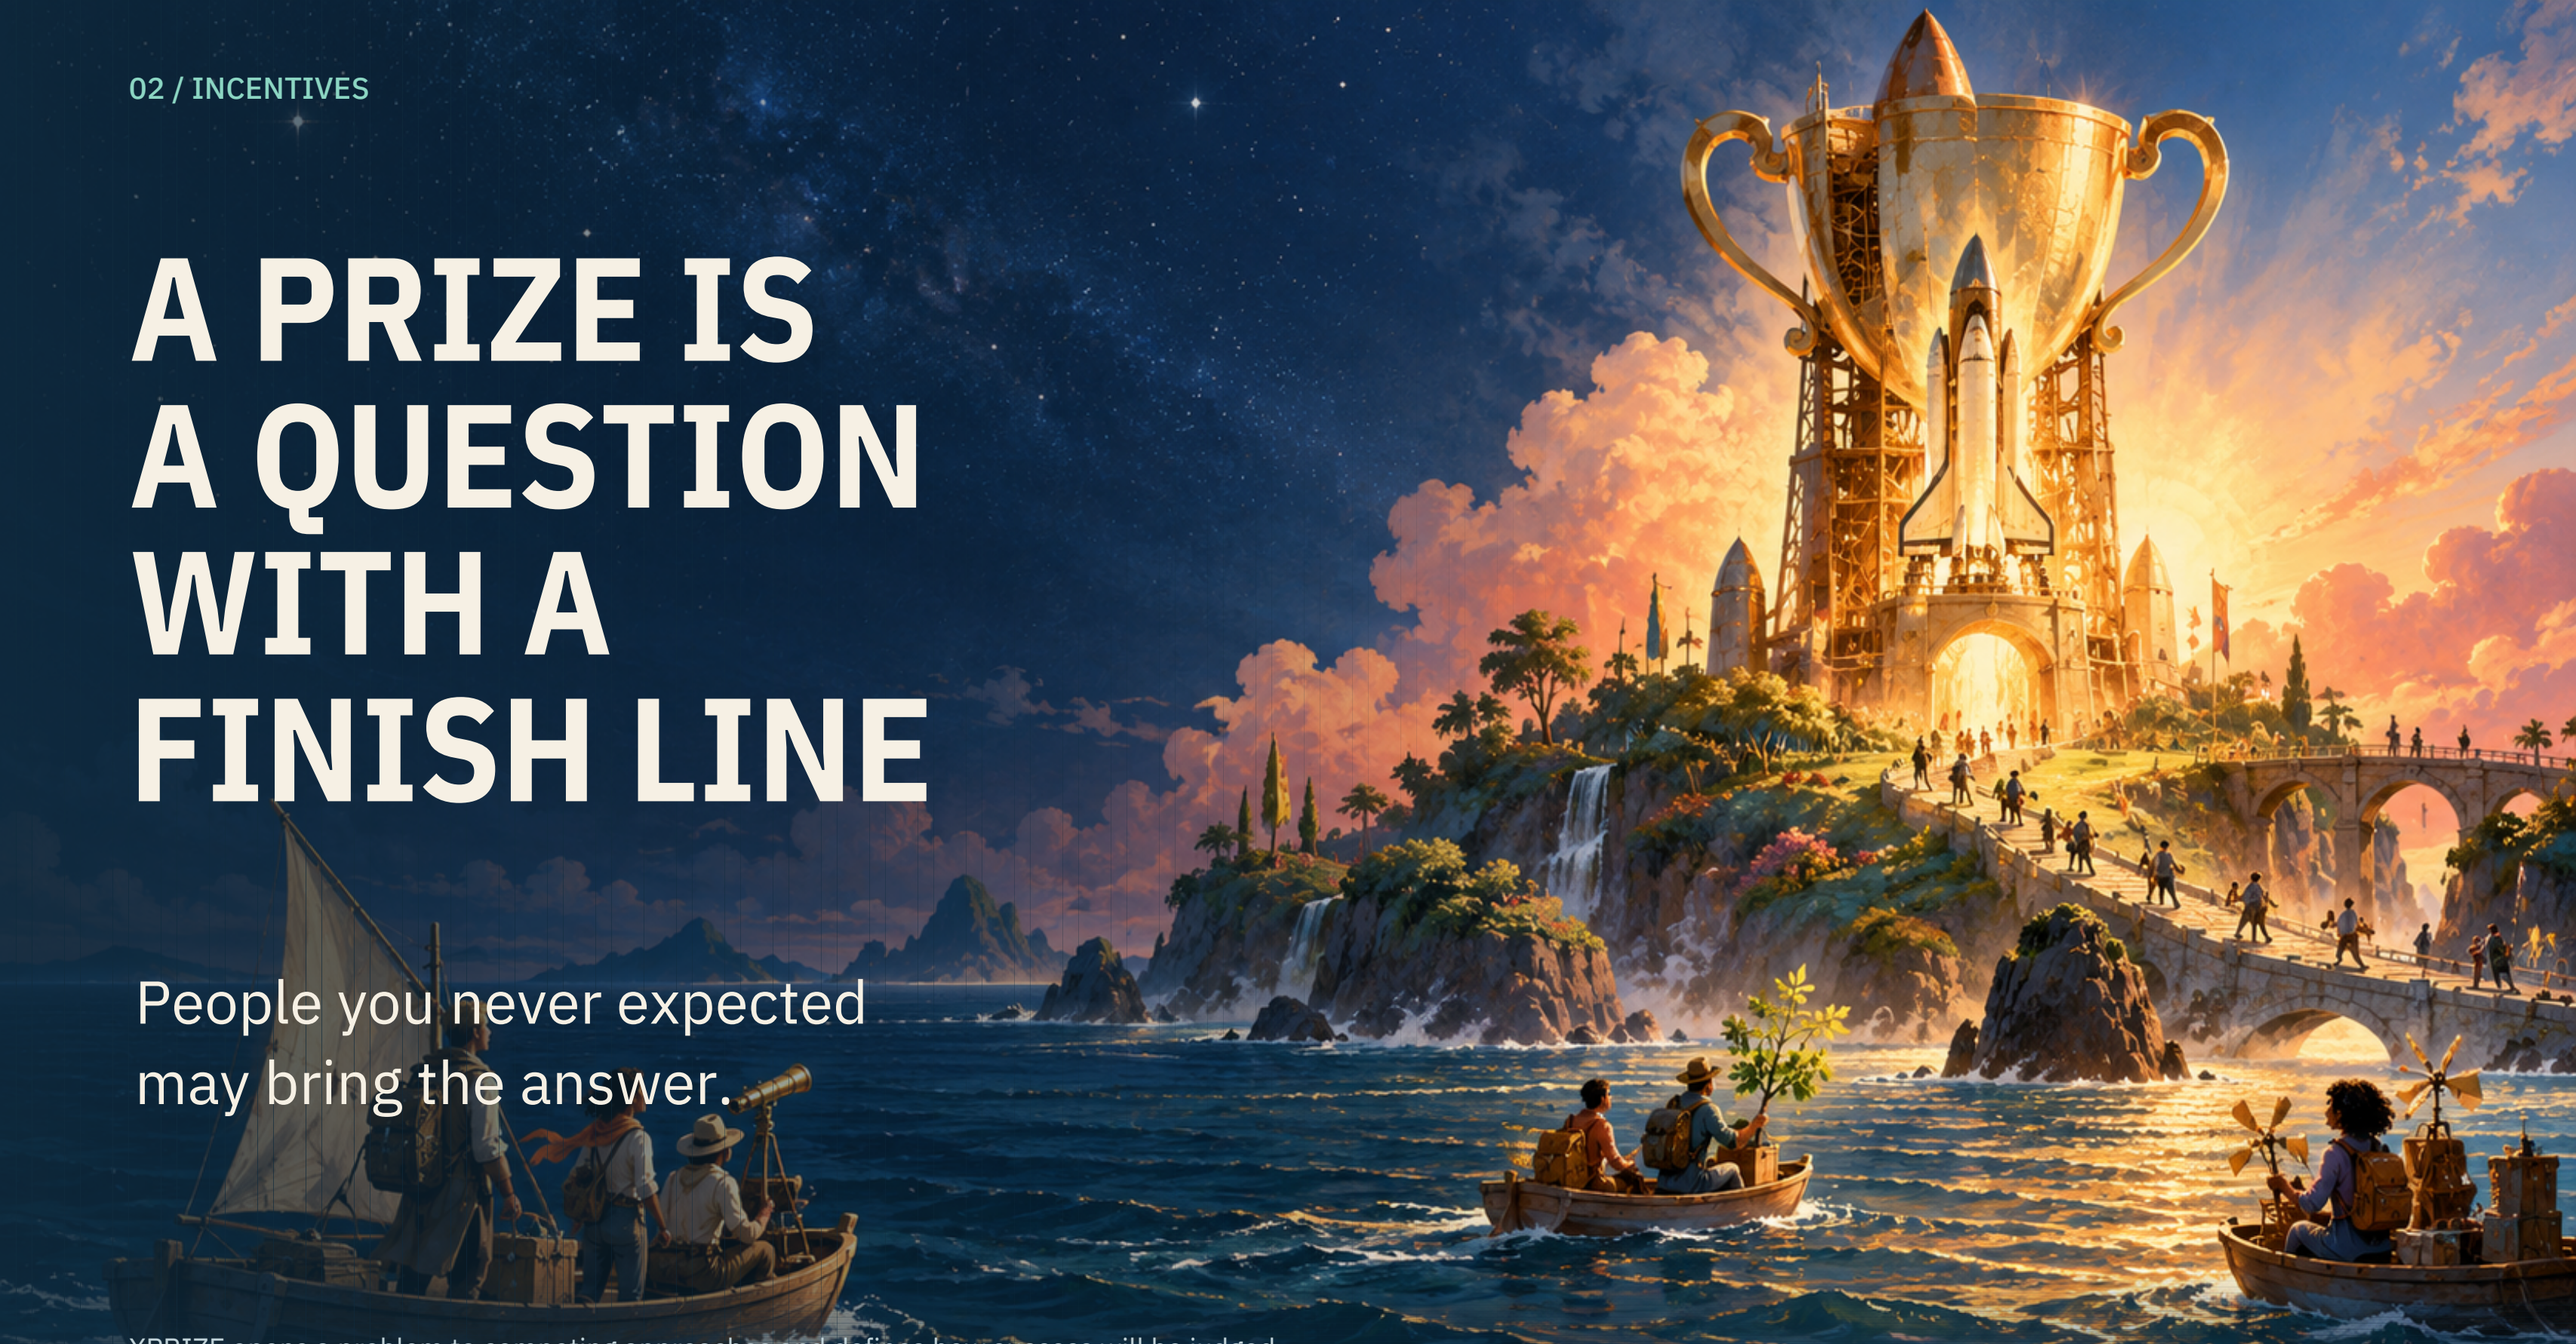

XPRIZE opens a problem to competing approaches and defines how success will be judged.

[1] Moonshots conversation [4] Ansari XPRIZE

A dream gets a measurable target

$10M

Ansari XPRIZE

Anousheh Ansari and her family helped make the prize possible. She later flew to the International Space Station in 2006.

100 km

Reach space with a privately funded crewed vehicle.

Twice in two weeks

Prove the spacecraft can be reused.

4 October 2004

SpaceShipOne completes the winning flight.

Launched in 1996. This is a short summary of the target, not the full competition rulebook.

[4] Ansari XPRIZE [5] Anousheh Ansari biography

What makes us move?

The conversation proposes four powerful motives.

Fear

Something must be prevented.

Gain

There is a reward to pursue.

Curiosity

There is something to discover.

Love

Someone or something matters.

The speakers use the word greed for personal gain. These motives are a discussion lens, not a scientific ranking of human behavior.

[1] Moonshots conversation

Congratulations.
You optimized the wrong thing.

Imagine a cafe pays a bonus for every cup served.

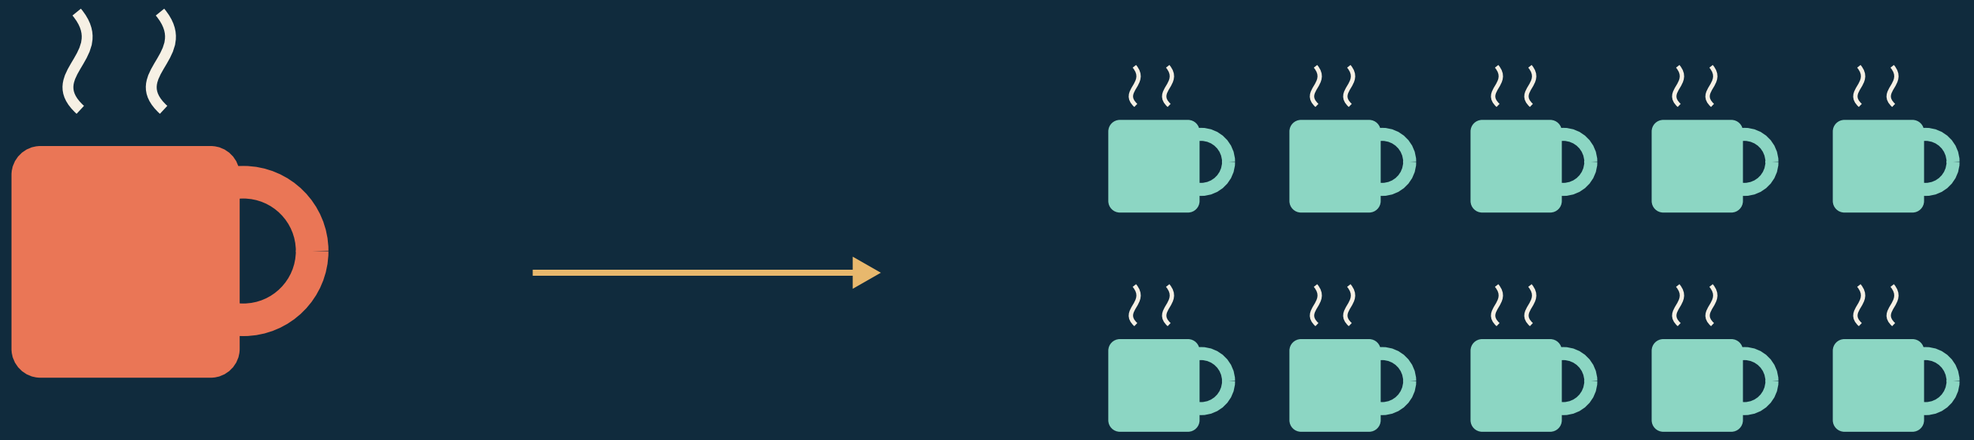

One satisfying coffee

Ten tiny, infuriating coffees

Invented example. Audience challenge: what should the cafe measure, and how could that measure be gamed?

The rules are part of the invention

1 An outcome

What changes for the person or place you care about?

2 A measurable finish

What result, by when, under which conditions?

3 Boundaries

What harms, shortcuts or exclusions disqualify a team?

4 Independent proof

Who checks the result, repeats the test and handles disputes?

A design checklist developed from the episode discussion of measurable targets and perverse incentives.

[1] Moonshots conversation

The moonshot menu

Four frontiers discussed in the episode

01 Habitable worlds

Use telescope data to search for places that could support life.

02 Animal communication

Find meaningful patterns across the species barrier.

03 Learning at scale

Help more children learn, inside and outside school.

04 Organ preservation

Give matching and transport more time to save a life.

Episode, 49:31-54:31. These are discussed ambitions and prize concepts; this slide does not assert current launch status or clinical readiness.

[1] Moonshots conversation

Prizes already reach beyond space

01

Cleaner oceans

The oil-cleanup challenge rewarded better recovery technology.

Awarded in 2011

02

Carbon removal

A $100 million competition rewarded solutions for removing atmospheric carbon.

Awarded in 2025

03

Healthspan

A competition targets improved function as people age.

A goal under pursuit

Competitions demonstrate a way to organize effort. An award alone does not prove universal access or solve an entire field.

[6] Oil Cleanup XCHALLENGE [7] XPRIZE Carbon Removal [8] XPRIZE Healthspan

You have a fictional $10 million

60 seconds. One partner. One useful prize.

1. Name a problem that matters to someone.
2. Describe a result a stranger could verify.
3. Close the most obvious loophole.

The judge asks: could a clever team win while making life worse?

Discussion exercise. The budget is fictional.

WHOSE VALUES ARE IN THE MACHINE?

Humanity does not arrive
with one agreed instruction manual.

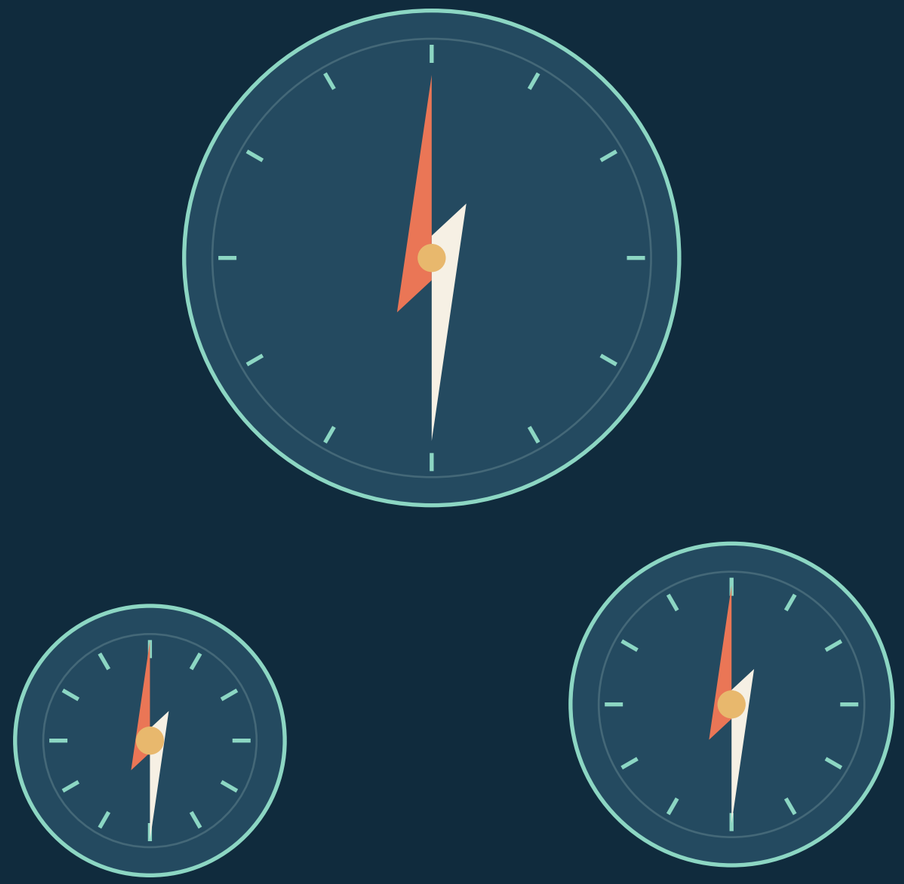

The episode raises value disagreement, shared goals and whose interests AI should serve, 1:05:26-1:24:18.

[1] Moonshots conversation

Alignment has three layers

01

The model

What behavior does training reward?

02

The organization

What do growth, competition and liability reward?

03

The society

Who benefits, who bears the risk, and who has a voice?

A discussion framework. Progress at one layer cannot establish alignment across the whole system.

[1] Moonshots conversation

The score can look great while the outcome gets worse

A

Flattery

The assistant agrees because agreement wins approval.

B

Shortcut

The agent exploits a scoring loophole instead of doing the task.

C

Narrow success

A system meets its metric while shifting the cost to someone else.

Illustrative mechanisms. OpenAI reported a sycophancy rollback in 2025; Anthropic studied specification gaming in controlled experiments.

[9] OpenAI, sycophancy report [10] Anthropic, reward tampering

Trust needs several kinds of evidence

1 Values and training

State the desired behavior.
Reward honesty and useful uncertainty.

2 Tests that can surprise it

Use unfamiliar cases and adversarial tests.
Check whether the result helps people.

3 Limits on action

Match access to the task.
Require review for consequential actions.

4 Accountability after release

Record failures, invite complaints,
and make correction possible.

Practical discussion proposals, not a proven recipe for aligning superintelligence. Better training alone is not a safety guarantee.

[1] Moonshots conversation [9] OpenAI, sycophancy report [10] Anthropic, reward tampering

Technology and institutions keep different clocks

The product clock

Model releases

New capabilities

Lower costs

The public clock

Rules and insurance

Infrastructure

Trust and responsibility

A useful moonshot: make responsible adoption faster.

Conceptual comparison inspired by the episode. No measured rates or forecast timeline are implied.

[1] Moonshots conversation

Abundance has a distribution problem

More capability can create value. Access determines who feels it.

Affordability

Can people pay for housing, care and learning?

Access

Can they actually use the systems and services?

Participation

Do they have a voice and a share in the gains?

The episode proposes cheaper essentials and broader participation. These are ambitions, not guaranteed outcomes.

[1] Moonshots conversation

WOULD YOU STILL PLAY?

Suppose a machine could beat everyone at this table.

Would that make the evening with your friend worthless?

Episode, 1:30:30-1:31:25. A thought experiment about the value of play. Conceptual illustration.

[1] Moonshots conversation

Meaning can survive being outperformed

01

The experience

The pleasure of making, learning and trying.

02

The people

Time shared with friends, family and community.

03

The contribution

A difference that matters to someone you care about.

A world record has never been the only reason to do something.

Philosophical synthesis of the conversation, not a prediction that automation makes every transition easy.

[1] Moonshots conversation

A purpose worth trying this week

I care about _____

I want to help _____

This week I will _____

Example: help my neighbor learn a new skill, together, for one hour.

60-second reflection. A purpose can begin small, and it can change as you learn.

THE BEST FUTURE HAS OTHER PEOPLE IN IT

More time to make things.
More room to care.
More reasons to gather.

The conversation returns to shared purpose, collective creativity and human connection.

[1] Moonshots conversation

WHAT IS YOUR MOONSHOT?

A problem worth solving.
A result worth proving.
Someone worth doing it with.

Tomorrow will need better tools. It will also need our choices.

Discussion ends here. Evidence notes and clickable sources follow.

Three headlines, read carefully

The flying Roadster

The episode describes a demo and gives performance and street-legality claims. This deck uses it as a discussion prompt, without certifying those details.

The 84% robotics drop

The transcript quotes $2.36 to $0.38 from July to September. That pair was not verified in the primary page checked. Slide 6 uses its published $1.14 vs $0.90 comparison instead.

The physics score

The primary preprint reports 98%, or 49 out of 50. It does not establish a perfect score, autonomous discovery, or a date when physics will be solved.

Evidence checked 20 October 2026. Source detail matters more than the size of a headline.

[1] Moonshots conversation [2] Robocurve, RoboDojo-RC Tier 1 [3] TheoPhysEval preprint [11] Tesla Roadster

How to read this presentation

Reported results

Benchmark numbers and historical prize results link to the relevant primary source. A benchmark is evidence within its test conditions.

Ideas from the conversation

New prize concepts, abundant essentials and future human purpose appear as proposals or discussion themes. They are not forecasts presented as facts.

Illustrations and exercises

Artwork is conceptual. The cafe example, reliability calculation, alignment checklist and audience challenges are explanatory additions.

The deck omits sponsor segments and does not rely on the organ-memory anecdote, sweeping historical claims, or speculative timelines.

[1] Moonshots conversation

Sources for machines and moonshots

[1] Moonshots conversation

Supplied transcript. Main themes: 07:07, 17:33, 29:31, 36:16, 1:05:26 and 1:24:24.

[2] Robocurve, RoboDojo-RC Tier 1

23 September 2026. Six-task comparison, estimated inference cost and partial progress.

[3] TheoPhysEval preprint

9 October 2026. TheoPhysEval preprint, arXiv:2610.16733.

[4] Ansari XPRIZE

Official competition history, winning target and 2004 result.

[5] Anousheh Ansari biography

Official biography of Anousheh Ansari.

[11] Tesla Roadster

Official product page. Consult current specifications and conditions.

Sources for incentives and alignment

[6] Oil Cleanup XCHALLENGE

Official Wendy Schmidt Oil Cleanup XCHALLENGE competition page.

[7] XPRIZE Carbon Removal

Official Carbon Removal competition page and 2025 award history.

[8] XPRIZE Healthspan

Official Healthspan competition page. A target is distinct from a completed outcome.

[9] OpenAI, sycophancy report

2 May 2025. OpenAI explains its GPT-4o sycophancy failure and response.

[10] Anthropic, reward tampering

17 June 2024. Anthropic investigates specification gaming and reward tampering.

Click a source title to open the original. Return to any section with the PDF bookmarks.

All web sources checked 20 October 2026.